

MASSACHUSETTS INSTITUTE OF TECHNOLOGY
DEPARTMENT OF MATHEMATICS
18.065/18.0651 - MATRIX METHODS (SPRING 2018)

GRADIENT DESCENT FOR LEAST SQUARES PROBLEMS

1. Gradient Descent

Recall that the least squares problem:

$$x_{ls} = \operatorname{argmin}_x \|Ax - b\|_2,$$

has the solution $x_{ls} = A^+b$ where A^+ is the pseudo-inverse. When the null-space of A is not empty, then x_{ls} is the minimum norm solution. When A is large, it can be computationally prohibitive to compute the SVD of A and then its pseudo-inverse before computing x_{ls} . In such settings, it can be shown that the iteration given by:

$$x_{k+1} = x_k - \mu A^H(Ax_k - b),$$

will minimize $\|Ax - b\|_2^2$ whenever $0 < \mu < 2/\sigma_1^2(A)$. Note that the iteration reaches a fixed point, i.e, $x_{k+1} = x_k$ when

$$A^H(Ax_k - b) = 0,$$

which are exactly the normal equations – so, the solution minimizes the least squares objective function! Your first task for this problem is to write a function called `lsgd` that implements the above least squares gradient descent algorithm. In Julia, your file should be named `lsgd.jl` and should contain the following function:

```
function lsgd(A, b, mu, x0, nIters)
#
# Syntax: x = lsgd(A, b, mu, x0, nIters)
#
# Inputs:  A is an m x n matrix
#          b is a vector of length m
#          mu is the step size to use, and must satisfy
#          0 < mu < 2 / norm(A)^2 to guarantee convergence
#          x0 is the initial starting vector (of length n) to use
#          nIters is the number of iterations to perform
#
# Outputs: x is a vector of length n containing the approximate solution
#
# Description: Performs gradient descent to solve the least squares problem
#
# \min_x \|b - A x\|_2
#
```

To check your solution, send your `lsgd.jl` file to `eeecs551@autograder.eecs.umich.edu`. Once your solution is correct, forward your confirmation email to `18065-code-submission@mit.edu` to receive credit for the lab.

After your code passes and you've forwarded your confirmation, use your solution to generate a plot of $\|x_{ls} - x_k\|$ as a function of k for A and b generated as:

```
m = 100; n = 50; sigma = 0.1
A = randn(m, n); xtrue = rand(n)
b = A * xtrue + sigma * randn(m);
```

Repeat the above experiment for $\sigma = 0.5, 1, 2$ and submit the plots with the problem set on Gradescope. Does $\|x_{ls} - x_k\|$ decrease monotonically with k in the plots?

2. Nesterov-Accelerated Gradient Descent

We now describe an accelerated gradient descent method due to Nesterov for the least squares problem:

$$x_{ls} = \operatorname{argmin}_x \|Ax - b\|_2,$$

that converges provably faster than the standard gradient descent method. The method consists of the following iteration:

$$\begin{aligned} t_{k+1} &= \frac{1 + \sqrt{1 + 4t_k^2}}{2} \\ z_{k+1} &= x_k + \frac{t_k - 1}{t_{k+1}}(x_k - x_{k-1}) \\ x_{k+1} &= z_{k+1} - \mu A^T(Az_{k+1} - b) \end{aligned}$$

initialized with $t_0 = 0$ and $x_{-1} = x_0$.

One can now show that the Nesterov-accelerated algorithm converges to x_{ls} when $0 < \mu < 1/\sigma_1^2(A)$. Your first task for this problem is to write a function called `lsngd` that implements the above Nesterov-accelerated least squares gradient descent algorithm.

In Julia, your file should be named `lsngd.jl` and should contain the following function:

```
function lsngd(A, b, mu, x0, nIters)
#
# Syntax: x = lsngd(A, b, mu, x0, nIters)
#
# Inputs:  A is an m x n matrix
#          b is a vector of length m
#          mu is the step size to use, and must satisfy
#          0 < mu < 1 / norm(A)^2 to guarantee convergence
#          x0 is the initial starting vector (of length n) to use
#          nIters is the number of iterations to perform
#
# Outputs: x is a vector of length n containing the approximate solution
#
# Description: Performs Nesterov-accelerated gradient descent to solve
#              the least squares problem
```

```
#
# \min x \|b - A x\|_2
#
```

To check your solution, send your lsngd.jl file to eeecs551@autograder.eecs.umich.edu. Once your solution is correct, forward your confirmation email to 18065-code-submission@mit.edu to receive credit for the lab.

After your code passes and you've forwarded your confirmation, use your solution to generate a plot of $\|x_{ls} - x_k\|$ as a function of k for A and b generated as:

```
m = 100; n = 50; sigma = 0.1
A = randn(m, n); xtrue = rand(n)
b = A * xtrue + sigma * randn(m);
```

Repeat the above experiment for $\sigma = 0.5, 1, 2$ and submit the plots with the problem set on Gradescope.

Compare the convergence characteristics for the accelerated gradient descent method to the standard gradient descent method. For a given step-size μ , which converges faster to the true solution? Does this hold for different values of (allowable) μ ? Turn in your plots illustrating the rate of convergence. Note that the convergence of the accelerated gradient descent method is not necessarily monotonic.